

pvdambtn7

January 3, 2023

# 1 SRM INSTITUTE OF SCIENCE AND TECHNOLOGY

## 1.1 MLSR - Regression Model

Project Done by : Bharath K M

[ ]:

## 2 1. Data Pre-processing

Import the required libraries

```
[3]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing

from warnings import filterwarnings
filterwarnings('ignore')
```

Load the CSV file

```
[4]: df_airline = pd.read_excel("airfare_CT3-1.xlsx")
df_airline.head(3)
```

```
[4]:
```

	Airline	Date	Departure Station	Arrival Station	\
0	IndiGo	24/03/2019	Banglore	New Delhi	
1	Air India	1/05/2019	Kolkata	Banglore	
2	Jet Airways	9/06/2019	Delhi	Cochin	

	Route Map	Departure Time	Arrival Time	Journey Time	Stops	\
0	BLR → DEL	22:20	01:10 22 Mar	2h 50m	non-stop	
1	CCU → IXR → BBI → BLR	05:50	13:15	7h 25m	2 stops	
2	DEL → LKO → BOM → COK	09:25	04:25 10 Jun	19h	2 stops	

	Extra Info	Price
0	No info	3897
1	No info	7662
2	No info	13882

```
[5]: df_airline.shape
```

```
[5]: (9000, 11)
```

```
[6]: df_airline.keys()
```

```
[6]: Index(['Airline', 'Date', 'Departure Station', 'Arrival Station', 'Route Map',
          'Departure Time', 'Arrival Time', 'Journey Time', 'Stops', 'Extra Info',
          'Price'],
          dtype='object')
```

```
[7]: df_airline.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9000 entries, 0 to 8999
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Airline                9000 non-null   object
1   Date                  9000 non-null   object
2   Departure Station      9000 non-null   object
3   Arrival Station        9000 non-null   object
4   Route Map              9000 non-null   object
5   Departure Time         9000 non-null   object
6   Arrival Time           9000 non-null   object
7   Journey Time           9000 non-null   object
8   Stops                 9000 non-null   object
9   Extra Info             9000 non-null   object
10  Price                  9000 non-null   int64
dtypes: int64(1), object(10)
memory usage: 773.6+ KB
```

### 2.0.1 Prepare the data

```
[8]: df_airline.describe()
```

```
[8]:
```

	Price
count	9000.000000
mean	9087.764333
std	4605.498942
min	1759.000000
25%	5228.000000

```
50%      8369.000000
75%     12373.000000
max      79512.000000
```

```
[9]: df_airline.dtypes
```

```
[9]: Airline      object
     Date        object
     Departure Station object
     Arrival Station object
     Route Map    object
     Departure Time object
     Arrival Time  object
     Journey Time  object
     Stops        object
     Extra Info    object
     Price         int64
     dtype: object
```

We can see from the above result that we have price as a numerical data with dtype int and the other data type as object.

```
[ ]:
```

### Perform missing value analysis

```
[10]: # sort the variables on the basis of total null values in the variable
      # 'isnull().sum()' returns the number of missing values in each variable

      missing_total = df_airline.isnull().sum()
      print(missing_total)
```

```
Airline      0
Date          0
Departure Station 0
Arrival Station 0
Route Map     0
Departure Time 0
Arrival Time  0
Journey Time  0
Stops         0
Extra Info    0
Price         0
dtype: int64
```

```
[ ]:
```

```
[ ]:
```

There are no missing values present in the given dataset

[ ]:

```
[11]: print(df_airline['Departure Station'].unique())
```

```
['Bangalore' 'Kolkata' 'Delhi' 'Chennai' 'Mumbai']
```

### 2.0.2 Replacing the departure values as per the station code

```
[12]: df_airline['Departure Station'] = df_airline['Departure Station'].
      ↪ replace({'Bangalore': 'BLR',
      ↪
      ↪ 'Delhi' : 'DEL',
      ↪
      ↪ 'Kolkata' : 'CCU',
      ↪
      ↪ 'Chennai' : 'MAA',
      ↪
      ↪ 'Mumbai' : 'BOM'})
      print(df_airline['Departure Station'].unique())
```

['BLR' 'CCU' 'DEL' 'MAA' 'BOM']

We have replaced the the departure station values as per the location code

## 2.1 Replacing arrival station values as per the location code

```
[13]: print(df_airline['Arrival Station'].unique())
```

```
['New Delhi' 'Bangalore' 'Cochin' 'Kolkata' 'Delhi' 'Hyderabad']
```

```
[14]: df_airline['Arrival Station'] = df_airline['Arrival Station'].
↳ replace({'Bangalore': 'BLR',
↳
↳ Delhi' : 'DEL',
↳
↳ : 'COK',
↳
↳ 'Hyderabad' : 'HYD',
↳
↳ 'Delhi' : 'DEL',
↳
↳ 'Kolkata' : 'CCU',
↳
↳ 'Chennai' : 'MAA',
```

```
↪ 'Mumbai' : 'BOM'})
print(df_airline['Arrival Station'].unique())
```

```
['DEL' 'BLR' 'COK' 'CCU' 'HYD']
```

```
[ ]:
```

```
[ ]:
```

### 2.1.1 Cleaning Extra info variable

```
[15]: print(df_airline['Extra Info'].unique())
```

```
['No info' 'In-flight meal not included' 'No check-in baggage included'
 '1 Short layover' 'No Info' '1 Long layover' 'Change airports'
 'Business class' 'Red-eye flight']
```

```
[16]: # There are two no info's {No info & No Info}.... Clearing those as the first
↪ step
```

```
df_airline['Extra Info'] = df_airline['Extra Info'].replace({"No info": "No
↪ Info"})
```

```
[17]: print(df_airline['Extra Info'].unique())
```

```
['No Info' 'In-flight meal not included' 'No check-in baggage included'
 '1 Short layover' '1 Long layover' 'Change airports' 'Business class'
 'Red-eye flight']
```

```
[18]: df_airline.groupby('Extra Info')['Extra Info'].count()
```

```
[18]: Extra Info
      1 Long layover          17
      1 Short layover          1
      Business class          3
      Change airports          4
      In-flight meal not included 1649
      No Info                7055
      No check-in baggage included 270
      Red-eye flight           1
      Name: Extra Info, dtype: int64
```

```
[19]: ## Assigning the categories using map function for the above categories shown
↪ in the result
```

```
df_airline['Extra Info'] = df_airline['Extra Info'].map({
```

```

    'No Info':0,
    'In-flight meal not included':1,
    'No check-in baggage included':2,
    '1 Long layover': 3,
    'Change airports':4,
    'Business class':5,
    '1 Short layover':6,
    'Red-eye flight':7
})

```

```
[20]: print(df_airline['Extra Info'].unique())
```

```
[0 1 2 6 3 4 5 7]
```

### 2.1.2 Cleaning Stops variabel

```
[21]: print(df_airline['Stops'].unique())
```

```
['non-stop' '2 stops' '1 stop' '3 stops']
```

```
[22]: df_airline['Stops'] = df_airline['Stops'].replace({'non-stop': 0,
                                                         '1 stop' : 1,
                                                         '2 stops': 2,
                                                         '3 stops':3})
```

```
[23]: print(df_airline['Stops'].unique())
```

```
[0 2 1 3]
```

### 2.1.3 Creating Day, Month, year variable from Date variable

```
[24]: df_airline.head(2)
```

```
[24]:
```

	Airline	Date	Departure Station	Arrival Station	\
0	IndiGo	24/03/2019	BLR	DEL	
1	Air India	1/05/2019	CCU	BLR	

	Route Map	Departure Time	Arrival Time	Journey Time	Stops	\
0	BLR → DEL	22:20	01:10 22 Mar	2h 50m	0	
1	CCU → IXR → BBI → BLR	05:50	13:15	7h 25m	2	

	Extra Info	Price
0	0	3897
1	0	7662

```
[25]: df_airline['Day'],df_airline['Month'],df_airline['Year'] = df_airline['Date'].
      ↪str.split('/',3).str
```

```
[26]: df_airline.head(2)
```

```
[26]:      Airline      Date Departure Station Arrival Station \
0      IndiGo  24/03/2019          BLR          DEL
1  Air India  1/05/2019          CCU          BLR

      Route Map Departure Time  Arrival Time Journey Time  Stops \
0          BLR → DEL        22:20  01:10 22 Mar        2h 50m      0
1  CCU → IXR → BBI → BLR        05:50        13:15        7h 25m      2

      Extra Info  Price Day Month  Year
0              0   3897  24    03  2019
1              0   7662   1    05  2019
```

#### 2.1.4 Dropping unwanted columns from the dataset

```
[27]: df_airline.drop('Date', axis='columns', inplace=True)
df_airline.drop('Arrival Time', axis='columns', inplace=True)
df_airline.drop('Year', axis='columns', inplace=True)
```

```
[28]: df_airline.head(2)
```

```
[28]:      Airline Departure Station Arrival Station      Route Map \
0      IndiGo          BLR          DEL          BLR → DEL
1  Air India          CCU          BLR  CCU → IXR → BBI → BLR

      Departure Time Journey Time  Stops  Extra Info  Price Day Month
0          22:20        2h 50m      0          0   3897  24    03
1          05:50        7h 25m      2          0   7662   1    05
```

```
[29]: df_airline.shape
```

```
[29]: (9000, 11)
```

## 2.2 3. Feature Engineering

### 2.2.1 Calculating distance

```
[30]: df_air_distance = pd.read_csv("air_distance.csv")
```

```
[31]: df_air_distance.head(2)
```

```
[31]:      Unnamed: 0 Source Dest  Distance(Km)
0              0    BLR  DEL        1709.71
1              1    CCU  IXR         327.84
```

```
[32]: import math

def getDistance(route):
    distance = 0.0
    route="".join(route.split())
    routeArray = route.split('→')
    i=0
    if len(routeArray) > 1:
        while i < (len(routeArray)-1):
            df_dist = df_air_distance[(df_air_distance['Source'] ==
→routeArray[i]) & (df_air_distance['Dest'] == routeArray[i+1])]
            if (df_dist.empty):
                df_dist = df_air_distance[(df_air_distance['Source'] ==
→routeArray[i+1]) & (df_air_distance['Dest'] == routeArray[i])]
            distValue = df_dist['Distance(Km)'].item()
            distance = distance + distValue
            i += 1
    return round(distance,2)
```

```
[33]: # df_airline['Distance(km)'] = distSeries.assign(distance = :
→getDistance(route))

df_airline['Distance(km)'] = df_airline['Route Map'].apply(lambda x:
→getDistance(x))

df_airline.head(3)
```

```
[33]:
```

	Airline	Departure Station	Arrival Station	Route Map	\
0	IndiGo	BLR	DEL	BLR → DEL	
1	Air India	CCU	BLR	CCU → IXR → BBI → BLR	
2	Jet Airways	DEL	COK	DEL → LKO → BOM → COK	

	Departure Time	Journey Time	Stops	Extra Info	Price	Day	Month	\
0	22:20	2h 50m	0	0	3897	24	03	
1	05:50	7h 25m	2	0	7662	1	05	
2	09:25	19h	2	0	13882	9	06	

	Distance(km)
0	1709.71
1	1838.55
2	2671.33



### 2.2.2 Creating arrival & departure hour, Minutes from arrival time and departure time

```
[34]: df_airline['Dep_Hr'],df_airline['Dep_Min'] = df_airline['Departure Time'].str.  
      ↪split(':',2).str  
df_airline['Duration'] = df_airline['Journey Time'].str.replace('h ','').str.  
      ↪replace('m','')  
df_airline['Duration_Hr'],df_airline['Duration_Min'] = df_airline['Duration'].  
      ↪str.split(':',2).str
```

### 2.2.3 Dropping unwanted columns

```
[35]: df_airline.drop('Departure Time', axis='columns', inplace=True)  
df_airline.drop('Journey Time', axis='columns', inplace=True)  
df_airline.drop('Duration', axis='columns', inplace=True)
```

```
[36]: df_airline.head(2)
```

```
[36]:
```

	Airline	Departure Station	Arrival Station	Route Map	Stops \
0	IndiGo	BLR	DEL	BLR → DEL	0
1	Air India	CCU	BLR	CCU → IXR → BBI → BLR	2

	Extra Info	Price	Day	Month	Distance(km)	Dep_Hr	Dep_Min	Duration_Hr \
0	0	3897	24	03	1709.71	22	20	2
1	0	7662	1	05	1838.55	05	50	7

	Duration_Min
0	50
1	25

```
[37]: df_airline.shape
```

```
[37]: (9000, 14)
```

### 2.2.4 Chnaging the datatype as per our requirment and model design

```
[38]: df_airline['Month'] = df_airline['Month'].astype(str).astype(int)  
df_airline.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 9000 entries, 0 to 8999  
Data columns (total 14 columns):  
#   Column                Non-Null Count  Dtype  
---  -  
0   Airline                9000 non-null   object  
1   Departure Station      9000 non-null   object  
2   Arrival Station        9000 non-null   object
```

```

3   Route Map          9000 non-null  object
4   Stops              9000 non-null  int64
5   Extra Info         9000 non-null  int64
6   Price              9000 non-null  int64
7   Day                9000 non-null  object
8   Month              9000 non-null  int32
9   Distance(km)       9000 non-null  float64
10  Dep_Hr             9000 non-null  object
11  Dep_Min            9000 non-null  object
12  Duration_Hr        9000 non-null  object
13  Duration_Min       8143 non-null  object
dtypes: float64(1), int32(1), int64(3), object(9)
memory usage: 949.3+ KB

```

```

[39]: # Replace the null values as 0

df_airline.Duration_Min.replace(np.nan, 0,inplace=True)

```

```

[40]: # String the duration_hr variable to remove unwanted spaces of special char

df_airline['Duration_Hr'] = df_airline['Duration_Hr'].str.rstrip('h')
df_airline.Duration_Hr.unique()

```

```

[40]: array(['2', '7', '19', '5', '4', '15', '21', '25', '13', '12', '26', '22',
        '23', '20', '10', '6', '11', '8', '16', '3', '27', '1', '14', '9',
        '18', '17', '24', '30', '28', '29', '37', '34', '38', '35', '36',
        '47', '33', '32', '31', '42', '39', '41'], dtype=object)

```

```

[41]: df_airline.head(2)

```

```

[41]:      Airline Departure Station Arrival Station      Route Map  Stops \
0      IndiGo          BLR          DEL      BLR → DEL          0
1  Air India          CCU          BLR  CCU → IXR → BBI → BLR      2

      Extra Info  Price Day  Month  Distance(km)  Dep_Hr  Dep_Min  Duration_Hr  \
0           0    3897   24     3      1709.71     22     20           2
1           0    7662    1     5      1838.55     05     50           7

      Duration_Min
0           50
1           25

```

```

[ ]:

```

## 2.3 4. Regularization

## 2.4 Renaming few variable to our understanding

```
[42]: df_airline = df_airline.rename(columns={'Departure Station': 'Source',
                                             'Arrival Station': 'Dest',
                                             "Extra Info": "Info"})
```

```
[43]: df_airline.head(2)
```

```
[43]:
```

	Airline	Source	Dest	Route	Map	Stops	Info	Price	Day	\
0	IndiGo	BLR	DEL	BLR → DEL		0	0	3897	24	
1	Air India	CCU	BLR	CCU → IXR → BBI → BLR		2	0	7662	1	

	Month	Distance(km)	Dep_Hr	Dep_Min	Duration_Hr	Duration_Min
0	3	1709.71	22	20	2	50
1	5	1838.55	05	50	7	25

### 2.4.1 Exporting the cleaned dataset as csv file

```
[44]: df_airline.to_csv('Cleaned_airline.csv', index=False)
```

## 3 2. Apply machine learning algorithm

```
[45]: from sklearn.preprocessing import LabelEncoder, MinMaxScaler
from sklearn.model_selection import train_test_split, GridSearchCV,
      ↪cross_val_score, RandomizedSearchCV
from sklearn.linear_model import LinearRegression, ElasticNet, Lasso, Ridge
from sklearn.metrics import mean_absolute_error, mean_squared_error
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor

import warnings
warnings.filterwarnings('ignore')

pd.set_option('display.max_columns', None)
```

```
[46]: df = pd.read_csv('Cleaned_airline.csv')
df.head(2)
```

```
[46]:
```


	Airline	Source	Dest	Route	Map	Stops	Info	Price	Day	\
0	IndiGo	BLR	DEL	BLR → DEL		0	0	3897	24	
1	Air India	CCU	BLR	CCU → IXR → BBI → BLR		2	0	7662	1	


	Month	Distance(km)	Dep_Hr	Dep_Min	Duration_Hr	Duration_Min
0	3	1709.71	22	20	2	50

1      5      1838.55      5      50      7      25


```
[47]: plt.figure(figsize=(12,6))
df.corr()['Price'].sort_values().plot(kind='bar');
```



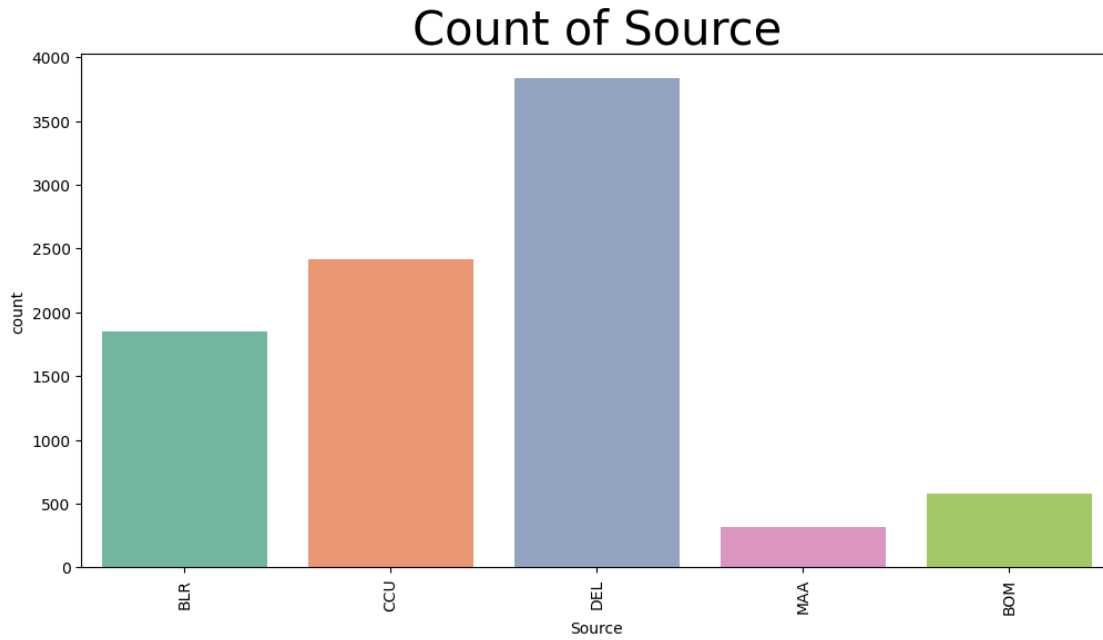
```
[48]: plt.figure(figsize=(12,6))
sns.countplot(x="Airline", data = df, palette='Set3')
plt.title('Count of Airlines', size=30)
plt.xticks(rotation=90)
plt.show()
```



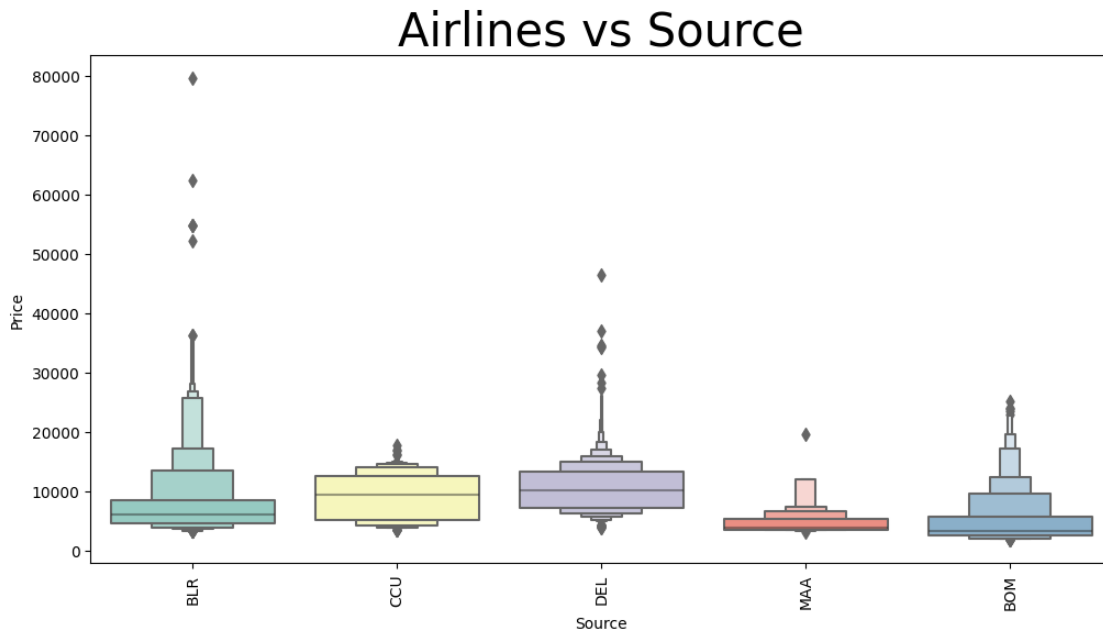
```
[49]: plt.figure(figsize=(12,6))
sns.boxenplot(x = 'Airline', y= 'Price', data=df, palette='Set3')
plt.title('Airlines vs Price', size=30)
plt.xticks(rotation=90)
plt.show()
```



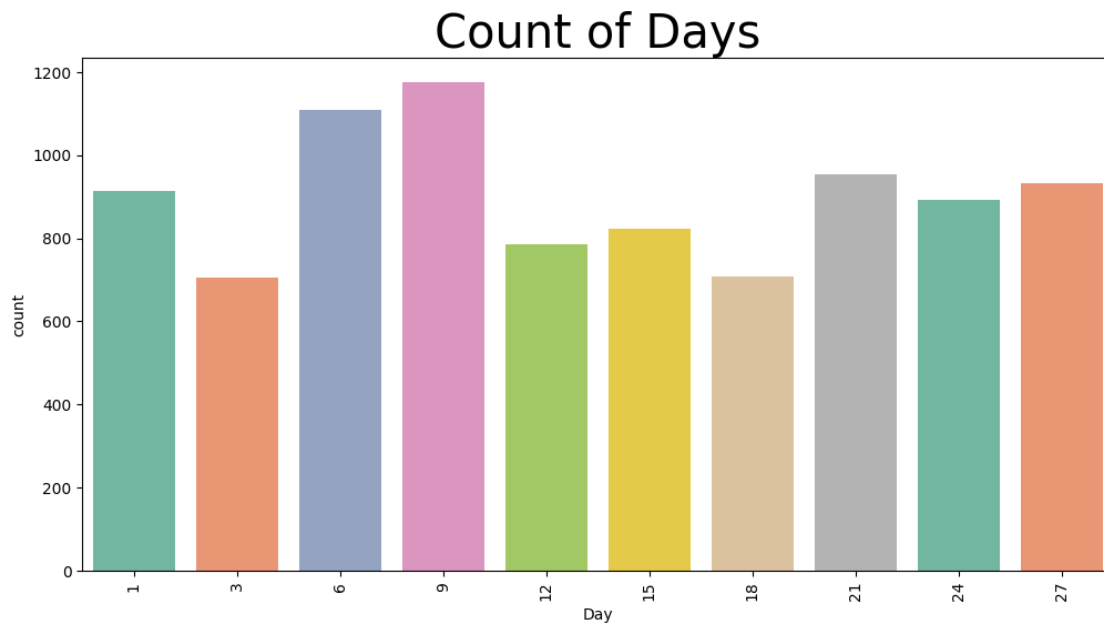
```
[50]: plt.figure(figsize=(12,6))
sns.countplot(x='Source', data = df, palette='Set2')
plt.title('Count of Source', size=30)
plt.xticks(rotation=90)
plt.show()
```



```
[51]: plt.figure(figsize=(12,6))
sns.boxenplot(x= 'Source', y= 'Price', data=df, palette='Set3')
plt.title('Airlines vs Source', size=30)
plt.xticks(rotation=90)
plt.show()
```

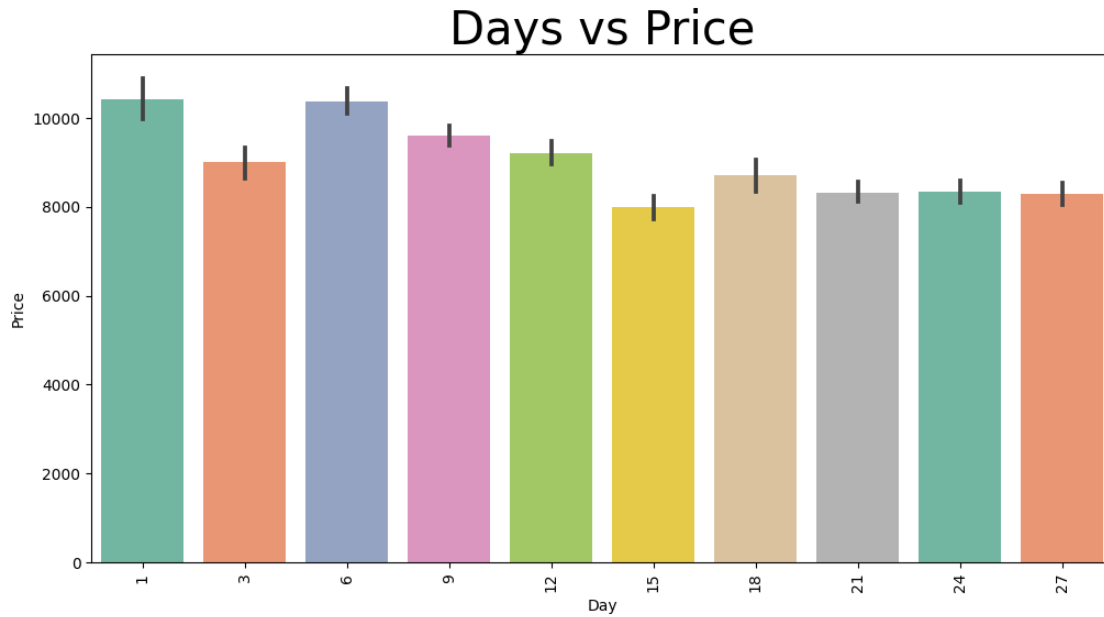


```
[52]: plt.figure(figsize=(12,6))
sns.countplot(x='Day', data= df, palette='Set2')
plt.title('Count of Days', size=30)
plt.xticks(rotation=90)
plt.show()
```




```
[53]: plt.figure(figsize=(12,6))
sns.barplot(x='Day', y='Price', data=df, palette='Set2')
plt.title('Days vs Price', size=30)
plt.xticks(rotation=90)
plt.show()
```






```
[54]: df['Month'] = df['Month'].map({  
      1: 'JAN',  
      2: 'FEB',  
      3: 'MAR',  
      4: 'APR',  
      5: 'MAY',  
      6: 'JUN',  
      7: 'JUL',  
      8: 'AUG',  
      9: 'SEP',  
     10: 'OCT',  
     11: 'NOV',  
     12: 'DEC'  
})
```

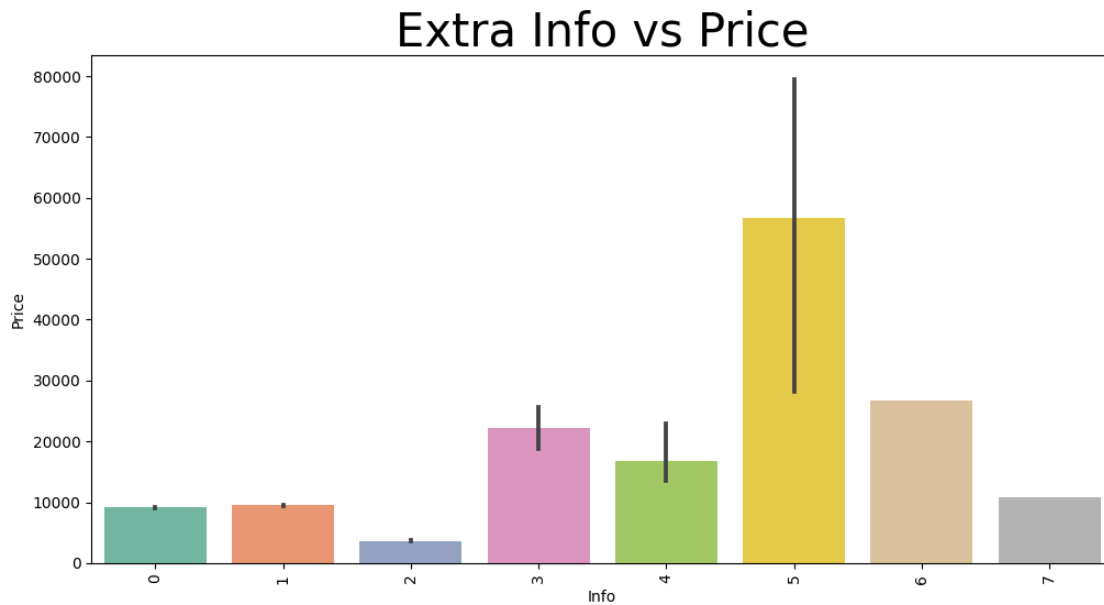
```
[55]: plt.figure(figsize=(12,6))  
sns.barplot(x='Month', y='Price', data=df, palette='Set2')  
plt.title('Month vs Price', size=30)  
plt.xticks(rotation=90)  
plt.show()
```




```
[56]: plt.figure(figsize=(12,6))
sns.barplot(x='Stops', y='Price', data=df, palette='Set2')
plt.title('Stops vs Price', size=30)
plt.xticks(rotation=90)
plt.show()
```



```
[57]: plt.figure(figsize=(12,6))
sns.barplot(x='Info', y='Price', data=df, palette='Set2')
plt.title('Extra Info vs Price', size=30)
plt.xticks(rotation=90)
plt.show()
```



```
[58]: df['Duration_bool'] = (df['Duration_Hr']*60)+df['Duration_Min']
plt.figure(figsize=(12,6))
sns.scatterplot(x= 'Duration_bool', y = 'Price', data=df, palette='Set2')
plt.title('Duration vs Price', size=30)
plt.xticks(rotation=90)
plt.show()
```



```
[59]: ncol=["Duration_bool"]
for i in ncol:
    q75, q25 = np.percentile(df.loc[:,i], [75 ,25])
    iqr = q75 - q25
    min = q25 - (iqr*1.5)
    max = q75 + (iqr*1.5)
    df = df.drop(df[df.loc[:,i] <= min].index)
    df = df.drop(df[df.loc[:,i] >= max].index)

df = df.dropna()
df1 = df[['Airline', 'Source', 'Dest', 'Stops',
          'Info', 'Price', 'Day', 'Month', 'Distance(km)', 'Duration_bool']]
df1 = df1.rename(columns={'Duration_bool': 'Duration'})
df1['Month'] = df1['Month'].map({
    'JAN':1,
    'FEB':2,
    'MAR':3,
    'APR':4,
    'MAY':5,
    'JUN':6,
    'JUL':7,
    'AUG':8,
    'SEP':9,
    'OCT':10,
    'NOV':11,
```

```

    'DEC':12
})
df.head(2)

```

```

[59]:
      Airline Source Dest      Route Map Stops Info Price Day \
0      IndiGo   BLR  DEL      BLR → DEL      0    0  3897  24
1  Air India   CCU  BLR  CCU → IXR → BBI → BLR      2    0  7662   1

      Month Distance(km) Dep_Hr Dep_Min Duration_Hr Duration_Min \
0      MAR      1709.71      22      20           2           50
1      MAY      1838.55       5      50           7           25

      Duration_bool
0           170
1           445

```

```

[60]: X = df1.drop('Price', axis=1)
      y = df1['Price']

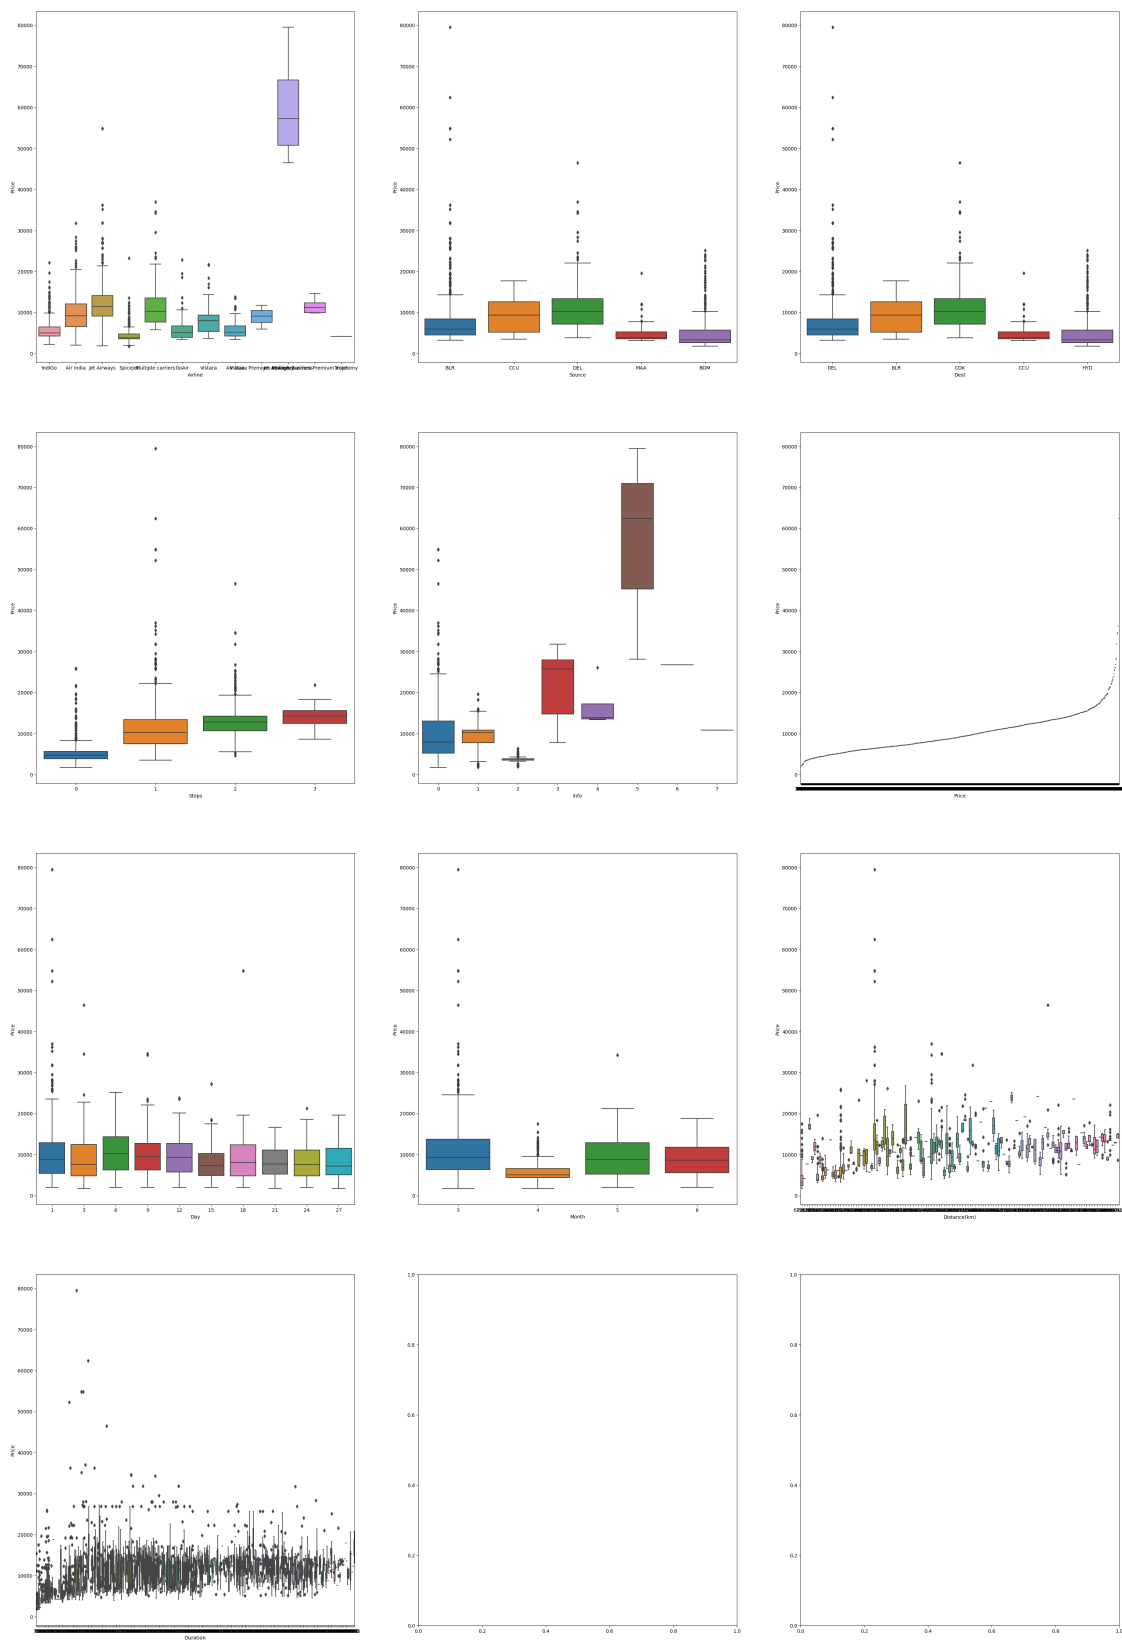
```

```

[61]: # set figure size
fig, ax = plt.subplots(4, 3, figsize=(40, 60))

# create box plot for categorical variables
for var, subplot in zip(df1.columns, ax.flatten()):
    sns.boxplot(x=var, y='Price', data=df1, ax=subplot)

```



```
[62]: df1.to_csv('final_airfare.csv', index=False)
```

```
[63]: # display all columns of the dataframe
pd.options.display.max_columns = None

# display all rows of the dataframe
pd.options.display.max_rows = None

# to display the float values upto 6 decimal places
pd.options.display.float_format = '{:.6f}'.format

# import train-test split
from sklearn.model_selection import train_test_split

# import various functions from statsmodels
import statsmodels
import statsmodels.api as sm

# import 'stats'
from scipy import stats

# 'metrics' from sklearn is used for evaluating the model performance
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error

# import function to perform linear regression
from sklearn.linear_model import LinearRegression

# import StandardScaler to perform scaling
from sklearn.preprocessing import StandardScaler

# import SGDRegressor from sklearn to perform linear regression with stochastic
↳ gradient descent
from sklearn.linear_model import SGDRegressor

# import function for ridge regression
from sklearn.linear_model import Ridge

# import function for lasso regression
from sklearn.linear_model import Lasso

# import function for elastic net regression
from sklearn.linear_model import ElasticNet

# import function to perform GridSearchCV
from sklearn.model_selection import GridSearchCV
```

```

from sklearn.ensemble import GradientBoostingRegressor

# import functions to perform cross validation
from sklearn.model_selection import LeaveOneOut
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import KFold

```

```

[64]: df = pd.read_csv('./final_airfare.csv')

# display first two observations using head()
df.head(2)

```

```

[64]:
      Airline Source Dest  Stops  Info  Price  Day  Month  Distance(km)  \
0    IndiGo   BLR  DEL      0     0   3897   24     3   1709.710000
1  Air India   CCU  BLR      2     0   7662    1     5   1838.550000

      Duration
0          170
1          445

```

```

[65]: # store the target variable 'PRICE' in a dataframe 'df_target'
df_target = df['Price']
df_feature = df.drop('Price', axis = 1)

df_num = df_feature.select_dtypes(include = [np.number])

# display numerical features
print("display numerical features:\n",df_num.columns)

df_cat = df_feature.select_dtypes(include = [np.object])

# display categorical features
print("display categorical features:\n",df_cat.columns)

# use 'get_dummies' from pandas to create dummy variables
# use 'drop_first' to create (n-1) dummy variables
dummy_var = pd.get_dummies(data = df_cat, drop_first = True)

```

```

display numerical features:
Index(['Stops', 'Info', 'Day', 'Month', 'Distance(km)', 'Duration'],
      dtype='object')
display categorical features:
Index(['Airline', 'Source', 'Dest'], dtype='object')

```

```

[66]: # initialize the standard scalar
X_scaler = StandardScaler()

```



```

# scale all the numeric variables
# standardize all the columns of the dataframe 'df_num'
num_scaled = X_scaler.fit_transform(df_num)

# create a dataframe of scaled numerical variables
# pass the required column names to the parameter 'columns'
df_num_scaled = pd.DataFrame(num_scaled, columns = df_num.columns)

# standardize the target variable explicitly and store it in a new variable 'y'
y = (df_target - df_target.mean()) / df_target.std()

```

```

[67]: # concat the dummy variables with numeric features to create a dataframe of all
      ↪ independent variables
# 'axis=1' concats the dataframes along columns
X = pd.concat([df_num_scaled, dummy_var], axis = 1)

# display first five observations
X.head(2)

```

```

[67]:
      Stops      Info      Day      Month  Distance(km)  Duration \
0 -1.221463 -0.479818  1.240175 -1.470566      -0.614115 -0.939403
1  1.789648 -0.479818 -1.474359  0.249940      -0.391266 -0.374705

      Airline_Air India  Airline_GoAir  Airline_IndiGo  Airline_Jet Airways \
0                      0              0              1              0
1                      1              0              0              0

      Airline_Jet Airways Business  Airline_Multiple carriers \
0                      0              0
1                      0              0

      Airline_Multiple carriers Premium economy  Airline_SpiceJet \
0                      0              0
1                      0              0

      Airline_Trujet  Airline_Vistara  Airline_Vistara Premium economy \
0                      0              0              0
1                      0              0              0

      Source_BOM  Source_CCU  Source_DEL  Source_MAA  Dest_CCU  Dest_COK \
0                0          0          0          0          0          0
1                0          1          0          0          0          0

      Dest_DEL  Dest_HYD
0            1          0
1            0          0

```

### 3.1 Train-test split

```
[68]: # split data into train subset and test subset
# set 'random_state' to generate the same dataset each time you run the code
# 'test_size' returns the proportion of data to be included in the testing set
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state = 10,
↳ test_size = 0.3)

# check the dimensions of the train & test subset using 'shape'
# print dimension of train set
print('X_train', X_train.shape)
print('y_train', y_train.shape)

# print dimension of test set
print('X_test', X_test.shape)
print('y_test', y_test.shape)
```

```
X_train (6254, 25)
y_train (6254,)
X_test (2681, 25)
y_test (2681,)
```

### 3.2 Creating RMSE values for train set

```
[69]: # create a generalized function to calculate the RMSE values for train set
def get_train_rmse(model):

    # For training set:
    # train_pred: prediction made by the model on the training dataset 'X_train'
    # y_train: actual values of the target variable for the train dataset

    # predict the output of the target variable from the train data
    train_pred = model.predict(X_train)

    # calculate the MSE using the "mean_squared_error" function

    # MSE for the train data
    mse_train = mean_squared_error(y_train, train_pred)

    # take the square root of the MSE to calculate the RMSE
    # round the value upto 4 digits using 'round()'
    rmse_train = round(np.sqrt(mse_train), 4)

    # return the training RMSE
    return(rmse_train)
```

### 3.3 Creating RMSE values for test data

```
[70]: # create a generalized function to calculate the RMSE values test set
def get_test_rmse(model):

    # For testing set:
    # test_pred: prediction made by the model on the test dataset 'X_test'
    # y_test: actual values of the target variable for the test dataset

    # predict the output of the target variable from the test data
    test_pred = model.predict(X_test)

    # MSE for the test data
    mse_test = mean_squared_error(y_test, test_pred)

    # take the square root of the MSE to calculate the RMSE
    # round the value upto 4 digits using 'round()'
    rmse_test = round(np.sqrt(mse_test), 4)

    # return the test RMSE
    return(rmse_test)
```

### 3.4 MAPE Calculation

```
[71]: # define a function to calculate MAPE
# pass the actual and predicted values as input to the function
# return the calculated MAPE
def mape(actual, predicted):
    return (np.mean(np.abs((actual - predicted) / actual)) * 100)

def get_test_mape(model):

    # For testing set:
    # test_pred: prediction made by the model on the test dataset 'X_test'
    # y_test: actual values of the target variable for the test dataset

    # predict the output of the target variable from the test data
    test_pred = model.predict(X_test)

    # calculate the mape using the "mape()" function created above
    # calculate the MAPE for the test data
    mape_test = mape(y_test, test_pred)

    # return the MAPE for the test set
    return(mape_test)
```

```
[ ]:
```

### 3.5 Creating a function to update scorecard

```
[72]: # create a function to update the score card for comparision of the scores from
      ↪different algorithms
      # pass the model name, model build, alpha and l1_ratio as input parameters
      # if 'alpha' and/or 'l1_ratio' is not specified, the function assigns '-'
      def update_score_card(algorithm_name, model, alpha = '-', l1_ratio = '-'):

          # assign 'score_card' as global variable
          global score_card

          # append the results to the dataframe 'score_card'
          # 'ignore_index = True' do not consider the index labels
          score_card = score_card.append({'Model_Name': algorithm_name,
                                          'Alpha (Wherever Required)': alpha,
                                          'l1-ratio': l1_ratio,
                                          'Test_MAPE': get_test_mape(model),
                                          'Test_RMSE': get_test_rmse(model),
                                          'R-Squared': get_score(model)[0],
                                          'Adj. R-Squared': get_score(model)[1]}, ignore_index =
      ↪True)
```

### 3.6 Function to plot barplot

```
[73]: # define a function to plot a barplot
      # pass the model
      def plot_coefficients(model, algorithm_name):
          # create a dataframe of variable names and their corresponding value of
      ↪coefficients obtained from model
          # 'columns' returns the column names of the dataframe 'X'
          # 'coef_' returns the coefficient of each variable
          df_coeff = pd.DataFrame({'Variable': X.columns, 'Coefficient': model.coef_})

          # sort the dataframe in descending order
          # 'sort_values' sorts the column based on the values
          # 'ascending = False' sorts the values in the descending order
          sorted_coeff = df_coeff.sort_values('Coefficient', ascending = False)

          # plot a bar plot with Coefficient on the x-axis and Variable names on
      ↪y-axis
          # pass the data to the parameter, 'sorted_coeff' to plot the barplot
          sns.barplot(x = "Coefficient", y = "Variable", data = sorted_coeff)

          # add x-axis label
          # set the size of the text using 'fontsize'
          plt.xlabel("Coefficients from {}".format(algorithm_name), fontsize = 15)
```

```

# add y-axis label
# set the size of the text using 'fontsize'
plt.ylabel('Features', fontsize = 15)

```

### 3.7 Function to generated R-squared and Adj R-squared

```

[74]: # define a function to get R-squared and adjusted R-squared value
def get_score(model):

```

```

    # score() returns the R-squared value
    r_sq = model.score(X_train, y_train)

    # calculate adjusted R-squared value
    # 'n' denotes number of observations in train set
    # 'shape[0]' returns number of rows
    n = X_train.shape[0]

    # 'k' denotes number of variables in train set
    # 'shape[1]' returns number of columns
    k = X_train.shape[1]

    # calculate adjusted R-squared using the formula
    r_sq_adj = 1 - ((1-r_sq)*(n-1)/(n-k-1))

    # return the R-squared and adjusted R-squared value
    return ([r_sq, r_sq_adj])

```

```

[75]: # n_splits: specify the number of k folds
kf = KFold(n_splits = 5)

```

```

[76]: # create a function 'get_score' that returns the R-squared score for the
      ↪ training set
# 'get_score' takes 5 input parameters
def Get_score(model, X_train_k, X_test_k, y_train_k, y_test_k):
    model.fit(X_train_k, y_train_k)                                # fit the
    ↪ model
    return model.score(X_test_k, y_test_k)

```

```

[77]: # create an empty list to store the scores
scores = []

# kf.split() splits the indices of X_train into train_index and test_index
# further dividing the X_train and y_train sets into train and test sets for
    ↪ cross validation
# Remember: Cross validation works on training set not on test set
# use '\\' for stacking the code

```

```

for train_index, test_index in kf.split(X_train):
    X_train_k, X_test_k, y_train_k, y_test_k = X_train.iloc[train_index], \
        X_train.iloc[test_index], \
        y_train.iloc[train_index], \
        y_train.iloc[test_index]

    # call the function 'get_scores()' and append the scores in the list
    scores.append(Get_score(LinearRegression(), X_train_k, X_test_k, y_train_k,
        y_test_k))

# print all scores
print('All scores: ', scores)

# print the minimum score from the list
# use 'round()' to round-off the minimum score upto 4 digits
# min() returns minimum score
print("\nMinimum score obtained: ", np.min(scores))

# print the maximum score from the list
# use 'round()' to round-off the maximum score upto 4 digits
# max() returns maximum score
print("Maximum score obtained: ", np.max(scores))

# print the average score from the list
# use 'round()' to round-off the average score upto 4 digits
# np.mean() returns average score
print("Average score obtained: ", np.mean(scores))

```

All scores: [0.583402449363021, 0.6215639618731148, 0.5779648965672155, 0.6145383837471873, -1.1033139815909515e+17]

Minimum score obtained: -1.1033139815909515e+17

Maximum score obtained: 0.6215639618731148

Average score obtained: -2.206627963181903e+16

```

[78]: # using cross_val_score() for k-fold cross validation
# estimator: pass the machine learning function. Here we are performing linear
# regression
# pass the X_train and y_train sets
# cv: stands for number of folds. Similar to k in KFold
# scoring: pass the scoring parameter e.g. 'r2' for r-squared,
# 'neg_mean_squared_error' for mean squared error (negative)
scores = cross_val_score(estimator = LinearRegression(),
                        X = X_train,
                        y = y_train,
                        cv = 5,

```

```
scoring = 'r2')
```

```
[79]: # print all scores
print('All scores: ', scores)

# print the minimum score from the list
# use 'round()' to round-off the minimum score upto 4 digits
# min() returns minimum score
print("\nMinimum score obtained: ", round(np.min(scores), 4))

# print the maximum score from the list
# use 'round()' to round-off the maximum score upto 4 digits
# max() returns maximum score
print("Maximum score obtained: ", round(np.max(scores), 4))

# print the average score from the list
# use 'round()' to round-off the average score upto 4 digits
# np.mean() returns average score
print("Average score obtained: ", round(np.mean(scores), 4))
```

```
All scores: [ 5.83402449e-01  6.21563962e-01  5.77964897e-01  6.14538384e-01
-1.10331398e+17]
```

```
Minimum score obtained: -1.1033139815909515e+17
```

```
Maximum score obtained: 0.6216
```

```
Average score obtained: -2.206627963181903e+16
```

```
[80]: # create an empty to store the MSE for each model
loocv_rmse = []

# instantiate the LOOCV method
loocv = LeaveOneOut()

# use the for loop to build the regression model for each cross validation
# use split() to split the dataset into two subsets; one with (n-1) data points
↳ and another with 1 data point
# where, n = total number of observations

for train_index, test_index in loocv.split(X_train):
    # create the train dataset, use iloc[] to retrieve the corresponding
    ↳ observations in train data
    # create the test dataset, use iloc[] to retrieve the corresponding
    ↳ observations in test data
    # use '\' for stacking the code
    X_train_l, X_test_l, y_train_l, y_test_l = X_train.iloc[train_index],
    ↳ X_train.iloc[test_index], \
```

```

y_train.iloc[train_index],
↪y_train.iloc[test_index]

# instantiate the regression model
linreg = LinearRegression()

# fit the model on training dataset
linreg.fit(X_train_1, y_train_1)

# calculate MSE using test dataset
# use predict() to predict the values of target variable
mse = mean_squared_error(y_test_1, linreg.predict(X_test_1))

# calculate the RMSE
rmse = np.sqrt(mse)

# use append() to add each RMSE to the list 'loocv_rmse'
loocv_rmse.append(rmse)

```

```

[81]: # print the minimum rmse from the list
# use 'round()' to round-off the minimum rmse upto 4 digits
# min() returns minimum rmse
print("\nMinimum rmse obtained: ", round(np.min(loocv_rmse), 4))

# print the maximum rmse from the list
# use 'round()' to round-off the maximum rmse upto 4 digits
# max() returns maximum rmse
print("Maximum rmse obtained: ", round(np.max(loocv_rmse), 4))

# print the average rmse from the list
# use 'round()' to round-off the average rmse upto 4 digits
# np.mean() returns average rmse
print("Average rmse obtained: ", round(np.mean(loocv_rmse), 4))

```

```

Minimum rmse obtained:  0.0001
Maximum rmse obtained:  3952236962.6924
Average rmse obtained:  631953.8728

```

```

[82]: models = [['LinearRegression', LinearRegression(), 'na'],
                ['ElasticNet', ElasticNet(), [{'alpha':[0.0001, 0.001, 0.01, 0.1, 1,
↪5, 10, 20, 40, 60],
                'l1_ratio':[0.0001, 0.0002, 0.001, 0.01, 0.1, 0.2]}]],
                ['Lasso', Lasso(), [{'alpha':[0.0001, 0.001, 0.01, 0.1, 1, 5, 10,
↪20]}]],
                ['Ridge', Ridge(), [{'alpha':[1e-4, 1e-3, 1e-2, 0.1, 1, 5, 10, 20, 40,
↪60, 80, 100]}]]],

```



```
['GradientBoostingRegressor', GradientBoostingRegressor(), 'na'],
['SGDRegressor', SGDRegressor(), 'na']]
```

```
[83]: # create an empty dataframe to store the scores for various algorithms
score_card = pd.DataFrame(columns=['Model_Name', 'Alpha (Wherever Required)',
    ↪ 'l1-ratio', 'R-Squared',
    ↪ 'Adj. R-Squared', 'Test_RMSE',
    ↪ 'Test_MAPE'])

for name, model, grid in models:
    model=model
    if grid == 'na':
        model.fit(X_train, y_train)
        update_score_card(algorithm_name = name, model = model)
    else:
        model = GridSearchCV(estimator = model,
                               param_grid = grid,
                               cv = 10)
        model.fit(X_train, y_train)
        update_score_card(algorithm_name = name, model = model, alpha = model.
    ↪ best_params_.get('alpha'),
                               l1_ratio = model.best_params_.get('l1_ratio'))
# sort the dataframe 'score_card' on 'Test_RMSE' in an ascending order using
    ↪ 'sort_values'
# 'reset_index' resets the index of the dataframe
# 'drop = True' drops the previous index
score_card = score_card.sort_values('Test_RMSE').reset_index(drop = True)

# color the cell in the column 'Test_RMSE' having minimum RMSE value
# 'style.highlight_min' assigns color to the minimum value
# pass specified color to the parameter, 'color'
# pass the data to limit the color assignment to the parameter, 'subset'
score_card.style.highlight_min(color = 'lightblue', subset = 'Test_RMSE')
```

```
[83]: <pandas.io.formats.style.Styler at 0x2a4038ecf70>
```

### 3.8 2. Module Creation

```
[84]: gradBoost = GradientBoostingRegressor()
gradBoost.fit(X_train, y_train)
prediction = gradBoost.predict(X_test)
print('RMSE : {}'.format(np.sqrt(mean_squared_error(y_test, prediction))))
```

```
RMSE : 0.3876990229902486
```


```
[85]: gradBoost.score(X_train, y_train), gradBoost.score(X_test, y_test)
```

```
[85]: (0.8340958488952346, 0.8513860966526166)
```

```
[86]: print('MAE:', mean_absolute_error(y_test, prediction))
      print('MSE:', mean_squared_error(y_test, prediction))
      print('RMSE:', np.sqrt(mean_squared_error(y_test, prediction)))
```

```
MAE: 0.2612347663124806
MSE: 0.15031053242759332
RMSE: 0.3876990229902486
```

```
[87]: plt.figure(figsize = (4,4))
      plt.scatter(y_test, prediction, alpha = 0.5)
      plt.xlabel("y_test")
      plt.ylabel("y_pred")
      plt.show()
```



## 4 Create a Pipeline and Save Predictive Model:

```
[88]: import pickle
      file = open('final_model.pkl', 'wb')
      pickle.dump(gradBoost, file)
```

```
[89]: model = open("final_model.pkl", "rb")  
      gradBoost = pickle.load(model)
```

```
[90]: from sklearn import metrics  
      predictions2=gradBoost.predict(X_test)  
      metrics.r2_score(y_test,predictions2)
```

```
[90]: 0.8513860966526166
```

So we have created predictive model and permanently saved in hard-drive with all the required pre-processing steps and whenever the new data to be tested

```
[ ]:
```